

Reinforcement Light Rays Path-Tracer Progress Report

Callum Pearce

cp15571@my.bristol.ac.uk

Abstract—This document contains my weekly progress of my thesis as part of my 4th year masters unit COMSM0111 in 2018/19. It is mainly useful to myself for reflecting on the decisions I have made throughout the project and why I chose them. However, I also hope it gives a detailed overview of how the project evolved with time for any other reader.

I. INTRODUCTION

Each section of this document describes a single week of work. For each section I have decided to include the following break down:

- **Goals:** A few set goals I aimed to achieve in that week and the motivation behind them.
- **Research/Implementation Details:** What was done, and how it was achieved.
- **Resources:** Describes what resources were notably helpful during the week for research/implementation details.
- **Reflection:** What I believed went well in the week, what to avoid in the future, and were the goal outlined achieved? Finally, if necessary, what has changed for the project as a whole?

WEEK 1

Building off some preliminary research I decided I had to build a ray-tracer in order to start any work with my project. This would be a good way of refreshing the basics of computer graphics.

A. Goals

- 1) Build a basic ray-tracer from scratch using only SDL and GLM as external libraries

B. Research/Implementation Goals

This week was fairly simple in terms of implementation. I mainly based my work on the ray-tracer I built with my project partner in my 3rd year of university. I used a similar project structure and followed the lab-sheets from COMS30015 by Carl Henrik Ek. By the end of the week I had built a ray-tracer which simulated the following in real-time:

- Constructing surfaces from triangle primitives and projecting them onto a 2D pixel plane for camera viewing via ray-tracing.
- Supports camera movement
- Direct illumination

C. Resources

As mentioned the main resources used were those provided for COMS30015 by Carl Henrik Ek in 2017. They gave me a good refresher on all the core concepts of ray-tracing and the mathematics behind it. I have based my rendering on the Cornell Box scene which is a classical scene for testing computer graphical renderings.

D. Reflection

I met my goal this week and built a well designed code-base to go with it.

WEEK 2

With a basic ray-tracer up and running, it was then time to add global illumination (indirect lighting) into the rendered scene and other features to make the basic ray-tracer complete for my purposes.

E. Goals

- 1) Implement Monte-Carlo global illumination within the ray-tracing pipeline (to go alongside direct light. It is Monte-Carlo global illumination I am planning to base my reinforcement learning technique presented by NVIDIA [1] on.
- 2) Create an object loader for the scene to test rendering in different scenes. I need a scene which has very low light levels in certain parts of the scene in order to show how reinforcement learning reduces noise in these areas.

F. Research/Implementation Goals

My implementation worked as follows; for every pixel in the image, calculate the colour by finding the direct light at that point combined with the indirect light. Where indirect light was sampled by shooting a ray into the scene and scattering it, then recursively finding the illumination at these points (via Monte Carlo).

For the object loader, I read in all vertices and then built triangles to form the defined surfaces in the file by using fan-triangulation. These triangles would be built into the scene by a call to the script with the file name of the .obj file to load.

I also introduced `openmp` to the project to speed things up by parallelising the ray-tracers pixel painting loop.

G. Resources

ScratchPixel 2.0 [4] provided an excellent description of Monte Carlo global illumination for a ray-tracer. As for the object loader, opengl gave a great tutorial for simple processing of .obj files [3].

H. Reflection

Monte Carlo global illumination was available for a custom scene as I had set out to achieve in this week. I also introduced an object loader which allows me to load in a scene of my choice, which will become especially useful when comparing different methods later on in the coursework. However, due to the introduction of global illumination the ray-tracer is far slower and takes a significant amount of time (nearly a day) to render a high quality image.

WEEK 3

With global illumination in hand and a custom scene to test it on the project moved on to begin working on reimplementing the *Learning Light transport the reinforced* way [1] paper.

I. Goals

- 1) Take detailed notes of the [1] paper and understand what needs to be changed with my current ray-tracer
- 2) Read through [5] Introduction, Markov Decision Processes and Temporal Difference learning chapters in order to understand the basis of Reinforcement learning.
- 3) Begin implementing the Radiance Volume data structure specified in [1].

J. Research/Implementation Goals

I began my week by first reading the Reinforcement Learning textbook [5] as to attain the second goal. I took many notes and gained a good understanding of temporal difference learning and how it builds on Markov Decision Processes as a model for reinforcement learning. I then read the NVIDIA paper [1] which I found out that I needed to use Expected Sarsa as my form of temporal difference learning for the path tracing algorithm created by NVIDIA.

From this reading I found that I needed to first convert my ray-tracer into a path-tracer, where we simulate light paths (rays) bouncing round the room until they hit a light surface. Only when the ray intersects with an area light does it gain luminance which then dictates the irradiance of a given point in the room along with the diffuse surfaces it bounced off. This required me to remodel my point light and convert it into an area light as well as redesign my implementation of `Triangles` and create the child classes `Surfaces` and `AreaLight` which extend from it. Whilst also modifying the ray-tracing Monte Carlo global illumination method to be a path-tracing Monte Carlo global illumination method [6]. Instead of scattering the rays every bounce, I instead sample many rays for a single pixel to begin with and just trace that single ray until it reaches a light source (or other terminating conditions).

With the path-tracer ready, I was able to begin working on reimplementing the data structure which [1] relies on, The Irradiance Volume [2]. The irradiance volume data structure stores the irradiance for many sampled points in the room, for which when you intersect with some geometry in the scene, you can interpolate between the precomputed irradiance values stored at these points to find the predicted irradiance at any point in the scene. The more sample points you have in the room the more accurate your estimate will be. For every sampled point the irradiance is calculated by finding the radiance at the given point by calculating all light incoming from uniformly sampled discretized angles (a grid converted into a hemisphere around the sample point). I managed to implement a single Irradiance Hemisphere and visualise it in the scene before the end of the week.

K. Resources

The most useful sources this week were the NVIDIA paper [1] (which I will no longer mention in these sections as the thesis is essentially based around it) and the reinforcement learning book [5].

L. Reflection

This week was different from those so far, I really focused on research and have a lot of notes to show for it, as compared to the past two weeks which have essentially been based around implementation of a ray-tracer. The work is getting more challenging and open to thought and interpretation of what cutting edge work has already been done by NVIDIA. I can safely say I met my 3 goals I outlined for this week.

WEEK 4

I now have a basis point to start off with the Irradiance Volume [2] implementation as I have successfully simulated a single `RadianceVolume` (class). It was now time to build this structure into the rendering pipeline and using these sampled point estimates into my approximation of global illumination for given point in the scene.

M. Goals

- 1) Find a way to sample radiance volumes around the scene uniformly and use their estimates in the rendering pipeline to create a perceptually realistic image.
- 2) Create KD-Tree in order to quickly look up the closest Radiance Volumes around a given point in the scene. Also introduce the `icc` Intel compiler to speed up programs performance.
- 3) Use Trilinear interpolation between the radiance volumes in the scene to find the irradiance estimate for a given point in the scene.
- 4) Begin implementing the reinforcement learning approach with the Irradiance Volume data structure.

N. Research/Implementation Goals

As for the first goal, I decided to sample the radiance volumes uniformly by sampling n radiance volumes on a given surface (using that surfaces normal), where the value of n is determined by the area of that triangle. This meant the algorithm is able to adapt to different scenes. This is different to what is proposed by [2] where a bilevel grid is used. This would likely give a more accurate estimate on surfaces as will focus all of our radiance volumes on surfaces rather than in the middle of the scene. This technique seems to produce good quality images.

The KD-Tree was a fairly straightforward implementation based upon code from my 3rd year computer graphics unit we implemented for a Photon Map. This made it quick to get working for an irradiance volume data structure and it made a massive improvement on query time, $O(\log n)$ (on average) to find the closest irradiance volume to given point in the scene.

With an efficient way to lookup the radiance estimate at a point in the scene, I created a way for visualizing the scene using interpolation between the stored radiance values alone. This is similar to a photon mapper and it enabled me to render images in almost real-time (after the pre-compute of the radiance volumes of course). I also implemented a barycentric coordinate system to interpolate between radiance volumes location to find an irradiance estimate. However, all of this was more or less just to check if my radiance map was implemented correctly by visualizing the result, it is not the rendering approach [1] take.

O. Resources

Clearly the Irradiance Volume paper [2] by Peter Shirley (who is an important person in graphics research) was a huge help for implementing the Radiance Volume, even if my version is quite heavily modified in how it used.

As for the KD-Tree my Github repository for my 3rd years constructed Photon Map was very useful.

P. Reflection

This week my work was heavily focused around creating this Radiance Map and visualizing it. I have managed to get some good evidence of its correctness and this is essential as it is key that it is implemented correctly for the reinforcement learning approach undertaken by [1] to work.

WEEK 5

With the radiance map created and validated, this week brings together a lot of what I have previously implemented and mainly involves implementing the reinforcement learning for light paths rendering approach.

Q. Goals

- 1) Implement an importance sampling approach based on the pre-computed data from the radiance map
- 2) Implement an initial version of the reinforcement learning approach and attempt to get similar results to that seen in [1]

R. Research/Implementation Goals

I quickly managed to get an importance sampling version of the path-tracer working by converting the radiance-volume values into a cumulative distribution which I could sample from the inverse of, in order to fire rays more likely towards light sources. This worked well at increasing the number of light rays which intersect with the area light source however some artefacts and strange patterns are present in the image, could this be to do with what [1] suggests about this not being a complete approximation?

I then moved on to implementing the reinforcement learning path tracer which was also fairly quick to implement. This does not perform the pre-computation step and clearly traces light paths towards the light far more effectively than randomly sampling, however once again there are artefacts present. I believe this is to do with my implementation of the radiance map, therefore I am going to plot the Voronoi render of my scene for the radiance volumes as shown in [1] to validate I have a similar layout. This will be effective in the future for debugging and visualizing the density of my samples. I have also realised that [1] might use a different sampling method for radiance volume locations using a low-discrepancy sequence (Hammersley) which I am planning on researching as well. I need to get this base datastructure correct before I can validate my reinforcement learning approach.

REFERENCES

- [1] Ken Dahn and Alexander Keller. Learning light transport the reinforced way. *arXiv preprint arXiv:1701.07403*, 2017.
- [2] Gene Greger, Peter Shirley, Philip M Hubbard, and Donald P Greenberg. The irradiance volume. *IEEE Computer Graphics and Applications*, 18(2):32–43, 1998.
- [3] opengl. Model loading. <http://www.opengl-tutorial.org/beginners-tutorials/tutorial-7-model-loading/>.
- [4] ScratchPixel. Scratchpixel2.0. <http://www.scratchapixel.com/>.
- [5] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. 2017.
- [6] Wikipedia. Path tracing. https://en.wikipedia.org/wiki/Path_tracing.